

# Datasets



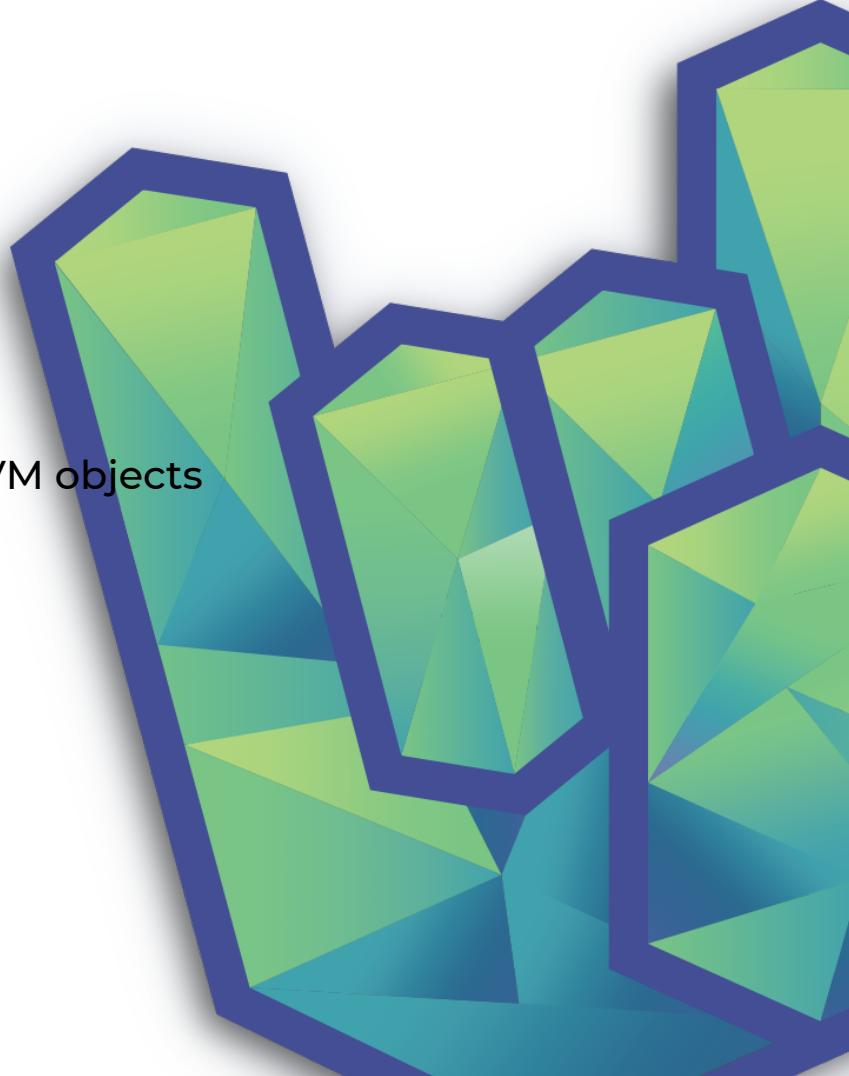
# Datasets, Part 2

Joining and grouping



# Datasets

Typed DataFrames: distributed collection of JVM objects



# Datasets

Distributed collection of JVM objects

Most useful when

- we want to maintain type information
- we want clean concise code
- our filters/transformations are hard to express in DF or SQL

Avoid when

- performance is critical: Spark can't optimize transformations

# Spark rocks

